

Lecture 16 — POSIX Threads (pthread)

Jeff Zarnett

jzarnett@uwaterloo.ca

Department of Electrical and Computer Engineering
University of Waterloo

March 24, 2024

The term `pthread` refers to the POSIX standard (also known as the IEEE 1003.1c standard) that defines thread behaviour in UNIX.

Wait, why are we talking about this?!

This is an embedded systems type course...

The pthread model is very common and used all across industry.

Knowing about this might seriously help you get a co-op job.

It will also help you if you want to take ECE 459 in the future...
Statistically, many of you do!

- `pthread_create`
- `pthread_exit`
- `pthread_join`
- `pthread_detach`
- `pthread_yield`
- `pthread_attr_init`
- `pthread_attr_destroy`
- `pthread_cancel`
- `pthread_testcancel`

```
pthread_create( pthread_t *thread,  
               const pthread_attr_t * attr,  
               void *(*start_routine)( void * ),  
               void *arg);
```

`thread`: a pointer to a pthread identifier and will be assigned a value when the thread is created.

`attr`: attributes; may be NULL for defaults.

`start_routine`: the function the new thread is to run.

`arg`: The argument passed to the routine we want to start.

The type of `start_routine` above is a function signature.

Thus, the `pthread_create` function has to be called with the name of a function matching that signature, such as:

```
void* do_something( void* start_params )
```

After creating a new thread, the process has two threads in it.

Scheduling of the threads is up to the operating system.

C: it is normal to have a single return value from a function, but usually we can have multiple input parameters.

But here we get only one of each?

Define a `struct` for the argument and return type!

```
void* function( void * void_arg ) {  
    parameters_t *arguments = (parameters_t*) args;  
    /* continue after this */  
}
```

We have to cast it inside the thread anyway...

The caller of the `pthread_create` function has to know what kind of argument is expected in the function being called.

Attributes can be used to set whether a thread is detached or joinable, scheduling policy, etc.

By default, new threads are usually joinable (that is to say, that some other thread can call `pthread_join` on them).

To prevent a thread from ever being joined, it can be created in the detached state (or use `pthread_detach`)

For virtually all scenarios that we will consider in this course the default values will be fine.

There is no mandatory hierarchy of threads.




Image Credit: Blaise Barney

New threads can create other threads.

The thread executes its function, until of course it gets to the end.

Usually, it will terminate with `pthread_exit`.

The use of `pthread_exit` is not the only way that a thread may be terminated.

Sometimes we want the thread to persist (hang around), but if we want to get a return value from the thread, then we need it to exit.

If a thread has no return values, it can just return `NULL`;

This will send `NULL` back to the thread that has joined it.

If the function that is called as a task returns normally rather than calling the exit routine, the thread will still be terminated.

Oh... Guess You Didn't Need This After All

Another way a thread might terminate is if the `pthread_cancel` function.
We'll come back to this topic in more detail soon.

A thread may also be terminated indirectly: if the entire process is terminated or if `main` finishes first (without calling `pthread_exit` itself).

End `main` with `pthread_exit` to automatically wait for all spawned threads.

Report, Number One!



Like the `wait` system call, the `pthread_join` is how we get a value out of the spawned thread:

```
pthread_join( pthread_t thread, void** retval );
```

`thread`: the thread you wish to join.

`retval`: wait... two stars?

Gotta Play the Level Again, Only Got 2 Stars

What we are looking for is a pointer to a void pointer.

That is, we are going to supply a pointer that the join function will update to be pointing to the value returned by that function.

Typically we supply the address of a pointer.

Maybe the example makes it clearer.

```
#include <stdlib.h>
#include <stdio.h>
#include <pthread.h>

void * run( void * argument ) {
    char* a = (char*) argument;
    printf("Provided_argument_is_%s!\n", a);
    int * return_val = malloc( sizeof( int ) );
    *return_val = 99;
    pthread_exit( return_val );
}

int main( int argc, char** argv ) {
    if (argc != 2) {
        printf("Invalid_args.\n");
        return -1;
    }
    pthread_t t;
    void* vr;

    pthread_create( &t, NULL, run, argv[1] );
    pthread_join( t, &vr );
    int* r = (int*) vr;
    printf("The_other_thread_returned_%d.\n", *r);
    free( vr );
    pthread_exit( 0 );
}
```


Next, we'll do an example where we don't use the return value of a thread, but do use attributes.

For the sake of simplicity: we are just going to count!



```
#include <pthread.h>
#include <stdio.h>

int sum; /* Shared Data */

void *runner(void *param);

int main( int argc, char **argv ) {
    pthread_t tid; /* the thread identifier */
    pthread_attr_t attr; /* set of thread attributes */

    if ( argc != 2 ) {
        fprintf(stderr, "usage: %s <integer_value>\n", argv[0]);
        return -1;
    }
    if ( atoi( argv[1] ) < 0 ) {
        fprintf(stderr, "%d must be >= 0\n", atoi(argv[1]));
        return -1;
    }

    pthread_attr_init( &attr ); /* set the default attributes */
    pthread_create( &tid, &attr, runner, argv[1] ); /* create the thread */
    pthread_join( tid, NULL );
    printf( "sum = %d\n", sum );
    pthread_attr_destroy( &attr );
    pthread_exit( NULL );
}
```

```
void *runner( void *param ) {  
    int upper = atoi( param );  
    sum = 0;  
    for ( int i = 1; i <= upper; i++ ) {  
        sum += i;  
    }  
    pthread_exit( 0 );  
}
```

In this example, both threads are sharing the global variable `sum`.

Do we have coordination?

Yes! The parent thread will join the newly-spawned thread (i.e., wait until it is finished) before it tries to print out the value.

If it did not, the parent would print the sum early.

Let's do a different take on that program.

When you watch Sesame Street and
can now count to 20.



My powers have doubled since the last time we met, Count.

```
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>

int sum = 0;

void* runner( void *param ) {
    int upper = atoi( param );
    for (int i = 1; i <= upper; i++ ) {
        sum += i;
    }
    pthread_exit( 0 );
}
```

```
int main( int argc, char** argv ) {

    pthread_t tid[3];

    if ( argc != 2 ) {
        printf("An integer value is required as an argument.\n");
        return -1;
    }
    if ( atoi( argv[1]) < 0 ) {
        printf( "%d must be >= 0.\n", atoi(argv[1]) );
    }

    for ( int i = 0; i < 3; ++i ) {
        pthread_create( &tid[i], NULL, runner, argv[1] );
    }
    for ( int j = 0; j < 3; ++j ) {
        pthread_join( tid[j], NULL );
    }
    printf( "sum = %d.\n", sum );

    pthread_exit( 0 );
}
```

What happens when we run this program?

For very small values of the argument, nothing goes wrong.

For a large number we get some strange and inconsistent results. Why?

There are three threads that are modifying sum.

Remember what we said about “at the same time”? We have to come back to it.

Thread cancellation is exactly what it sounds like: a running thread will be terminated before it has finished its work.

The thread that we are going to cancel is called the *target*.



1 Asynchronous Cancellation

2 Deferred Cancellation

thread can declare its own cancellation type through the use of the function:

```
pthread_setcanceltype( int type, int *oldtype )
```

type: PTHREAD_CANCEL_DEFERRED or PTHREAD_CANCEL_ASYNCHRONOUS

oldtype: previous state, if we care.

The pthread command to cancel a thread is `pthread_cancel` and it takes one parameter (the thread identifier).

To check if the current thread has been cancelled, the function call is `pthread_testcancel` which takes no parameters.


It's polite to check this, if it's a risk.

A large number of functions are **cancellation points**.

That is, the POSIX specification requires there is an implicit check for cancellation when calling one of those functions.

Even more are “potential cancellation points” – maybe, maybe not?

Sometimes a thread could die before it has cleaned up.



This can leave memory allocated, things locked...

We can prevent this with cancellation handlers.

The functions for cleaning up are:

```
/* Register cleanup handler, with argument */  
pthread_cleanup_push( void (*routine)(void*), void *argument );  
/* Run if execute is non-zero */  
pthread_cleanup_pop( int execute );
```

The push function always needs to be paired with the pop function at the same level in your program (where level is defined by the curly braces).

Consider the following code:

```
void* do_work( void* argument ) {  
    struct job * j = malloc( sizeof( struct job ) );  
    /* Do something useful with this structure */  
    /* Actual work to do not shown */  
    free( j );  
    pthread_exit( NULL );  
}
```

```
void cleanup( void* mem ) {
    free( mem );
}

void* do_work( void* argument ) {
    struct job * j = malloc( sizeof( struct job ) );
    pthread_cleanup_push( cleanup, j );
    /* Do something useful with this structure */
    /* Actual work to do not shown */
    free( j );
    pthread_cleanup_pop( 0 ); /* Don't run */
    pthread_exit( NULL );
}
```
